

Case Study:


Applying Exploratory Data Analysis Techniques to the Titanic Dataset Using R



By:

Amr Eleraqi

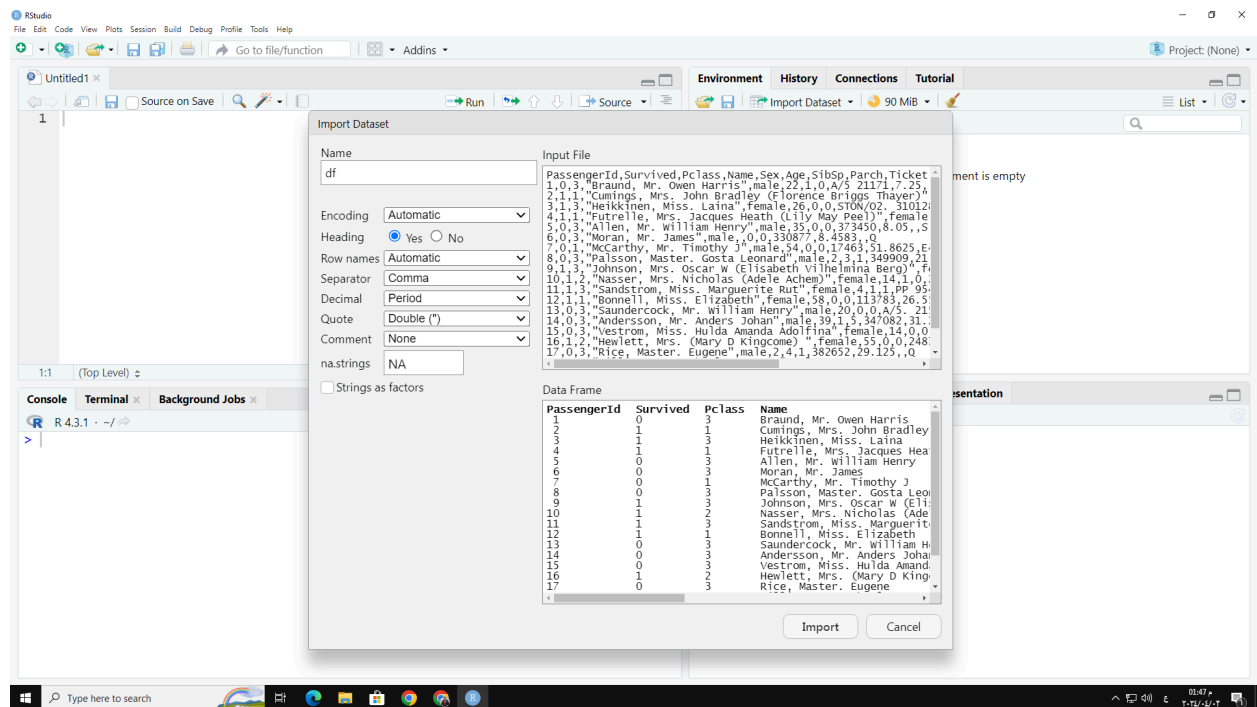
aeleraqi@aucegypt.edu



 **Data Source:** <https://github.com/datasciencedojo/datasets/blob/master/titanic.csv>

Importing Data:

My preferred approach to import data into R Studio is through the use of the "Import Data" option found within the Environment tab. This feature offers an array of customizable settings tailored to the specific type of file being imported, whether it be a CSV or XLS document.



Exploring the Data:

After successfully importing the data set into R Studio, there are multiple functions available to help us explore the content effectively. Among those, `View()`, `head()` and `tail()` are particularly useful when navigating through large datasets. Let me briefly describe each of them below:

- `View()`: With `view()`, you can display the entire data frame interactively, enabling easy navigation and exploration of individual cells, columns, and rows directly within RStudio.
- `head()`: When dealing with extensive data sets, sometimes it is more practical to examine just a few initial lines rather than viewing the whole table. That's where `head()` comes in handy. It returns the first `n` observations (rows) of a given data frame, defaulting to six if no argument is provided.

```
head(df, n = 10) # Display the first 10 rows instead
```
- `tail()`: Similar to `head()`, `tail()` displays the last `n` observations of the data frame, showing six

rows by default.

```
tail(data_frame_name, n = 5) # Display only the last 5 rows
```

- `dim()`: returns the dimensions of the data frame, specifically indicating how many rows and columns it has.

Understanding the measurements:

After obtaining a general overview of the data frame using the above functions, delving deeper becomes crucial to thoroughly comprehend the nature of the variables and their respective measurements present in the dataset. Understanding the exact meaning of each column and knowing the measurement scales facilitate accurate downstream data analysis and interpretation.

Here are several key functions in R that allow us to acquire detailed insights into your data frame:

names()

Returns character vectors representing the column names of a data frame

```
> names(df)
[1] "PassengerId" "Survived"    "Pclass"      "Name"        "Sex"
[6] "Age"          "SibSp"       "Parch"       "Ticket"      "Fare"
[11] "Cabin"        "Embarked"
```

str()

Displays the internal structure of objects, providing details on class, length, and modes of components

```
> str(df)
'data.frame':      891 obs. of  12 variables:
 $ PassengerId: int   1  2  3  4  5  6  7  8  9 10 ...
 $ Survived   : int   0  1  1  1  0  0  0  0  1  1 ...
 $ Pclass     : int   3  1  3  1  3  3  1  3  3  2 ...
 $ Name       : chr   "Braund, Mr. Owen Harris" "Cumings, Mrs. John Bradley
(Florence Briggs Thayer)" "Heikkinen, Miss. Laina" "Futrelle, Mrs. Jacques
Heath (Lily May Peel)" ...
 $ Sex        : chr   "male" "female" "female" "female" ...
 $ Age        : num   22  38  26  35  35 NA  54  2  27  14 ...
 $ SibSp      : int   1  1  0  1  0  0  0  3  0  1 ...
 $ Parch      : int   0  0  0  0  0  0  0  1  2  0 ...
 $ Ticket     : chr   "A/5 21171" "PC 17599" "STON/O2. 3101282" "113803" ...
 $ Fare       : num   7.25 71.28 7.92 53.1 8.05 ...
 $ Cabin      : chr   "" "C85" "" "C123" ...
 $ Embarked   : chr   "S" "C" "S" "S" ...
```

Brief descriptions for each variable:

Name	Title	Description
PassengerId	Passenger Identifier	Unique ID assigned to each passenger
Survived	Survivorship Status	Flag indicating whether the passenger lived (1) or died (0)
Pclass	Ticket Fare Class	Travel class category: First, Second, Third
Name	Passenger Full Name	Legal name of the passenger
Sex	Passenger Gender	Male or Female gender classification
Age	Age at Embarkation	Continuous age variable measured in years
SibSp	Family Members	Quantitative metric for sibling and spouse presence
Parch	Parents and Children	Counter measuring parental units and offspring
Ticket	Booking Reference	Alpha-numeric reference string tied to the ticket
Fare	Voyage Expense	Cost paid for the journey expressed in monetary terms
Cabin	Accommodation Space Assignment	Label designating allocated cabin compartment
Embarked	Boarding Location	Departure port of the passenger

Handling the Missing Values:**1. Calculate the number of NA values in each column of df**

```
> colSums(is.na(df))
```

PassengerId	Survived	Pclass	Name	Sex	Age
0	0	0	0	0	177
SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	0	0	0	0	0

2. Remove rows with missing values

```
clean_df <- na.omit(df)
```

3. Replacing the missing values with the mean

The following command replaces any missing values in the Age column of the data frame df with the mean of the non-missing values in the same column

```
> df$Age[is.na(df$Age)] <- mean(df$Age, na.rm = T)
```



Replacing missing values with the mean or median depends upon the underlying assumptions, level of tolerance towards extreme values, and sensitivity to influential data points. Both approaches have pros and cons, which necessitate careful consideration before applying them to fill in gaps caused by missingness.

This command counts the number of missing values in the Age column

```
> sum(is.na(df$Age))  
[1] 0
```

calculates the mean of the non-missing values in the Age column

```
> mean(df$Age, na.rm = T)  
[1] 29.69912
```

calculates the mean of the Age column

```
> mean(df$Age)  
[1] 29.69912
```

Transforming Selected Variables into Factors

`as.factor()` conversions convert the Survived, Pclass, and Sex columns to factors. Recall that factors represent categorical variables encoded as integer values internally accompanied by labeled levels. This command converts the Survived, Pclass and Sex columns of the data frame `df` into factors

```
> df$Survived <- as.factor(df$Survived)  
> df$Pclass <- as.factor(df$Pclass)  
> df$Sex <- as.factor(df$Sex)
```

The following command checks if the Survived, Pclass and Sex columns of the data frame `df` is a factor


```
> is.factor(df$Survived)  
[1] TRUE
```

```
> levels(df$Survived)  
[1] "0" "1"
```

One alternative method involves employing the `ifelse` function combined with assignment operators to create new factor variables with desired labels. Here's the modified example:

```
df$Survived <- ifelse(df$Survived==1, "Yes", "Yes")  
  
> levels(df$Survived)  
[1] "Yes" "Yes"
```


```
> str(df)
'data.frame':      891 obs. of  12 variables:
 $ PassengerId: int  1 2 3 4 5 6 7 8 9 10 ...
 $ Survived   : Factor w/ 2 levels "0","1": 1 2 2 2 1 1 1 1 2 2 ...
 $ Pclass     : Factor w/ 3 levels "1","2","3": 3 1 3 1 3 3 1 3 3 2 ...
 $ Name       : chr  "Braund, Mr. Owen Harris" "Cumings, Mrs. John Bradley
(Florence Briggs Thayer)" "Heikkinen, Miss. Laina" "Futrelle, Mrs. Jacques
Heath (Lily May Peel)" ...
 $ Sex        : Factor w/ 2 levels "female","male": 2 1 1 1 2 2 2 2 1 1 ...
 $ Age        : num  22 38 26 35 35 NA 54 2 27 14 ...
 $ SibSp      : int  1 1 0 1 0 0 0 3 0 1 ...
 $ Parch      : int  0 0 0 0 0 0 0 1 2 0 ...
 $ Ticket     : chr  "A/5 21171" "PC 17599" "STON/O2. 3101282" "113803" ...
 $ Fare       : num  7.25 71.28 7.92 53.1 8.05 ...
 $ Cabin      : chr  "" "C85" "" "C123" ...
 $ Embarked   : chr  "S" "C" "S" "S" ...
```

 For discrete variables transformed into factors, the `str(df)` command would reveal the shift from their original data types to factors, along with associated factor level information.

```
> summary(df)
 PassengerId   Survived  Pclass         Name             Sex
Min.   :  1.0   0:549    1:216   Length:891   female:314
1st Qu.:223.5   1:342    2:184   Class :character  male  :577
Median :446.0           3:491   Mode  :character
Mean   :446.0
3rd Qu.:668.5
Max.   :891.0

      Age      SibSp      Parch      Ticket
Min.   : 0.42   Min.   :0.000   Min.   :0.0000   Length:891
1st Qu.:20.12   1st Qu.:0.000   1st Qu.:0.0000   Class :character
Median :28.00   Median :0.000   Median :0.0000   Mode  :character
Mean   :29.70   Mean   :0.523   Mean   :0.3816
3rd Qu.:38.00   3rd Qu.:1.000   3rd Qu.:0.0000
Max.   :80.00   Max.   :8.000   Max.   :6.0000
NA's   :177

      Fare      Cabin      Embarked
Min.   :  0.00   Length:891   Length:891
1st Qu.:  7.91   Class :character  Class :character
Median : 14.45   Mode  :character  Mode  :character
Mean   : 32.20
3rd Qu.: 31.00
Max.   :512.33
```

 Regarding the output generated by `summary(df)`, noticeable modifications arise in the presentation of categorical variables. Instead of reporting statistical measures typically used for continuous data, the resulting table shows frequency counts across factor categories.

Subsetting Data Frames

1. Selecting Specific Columns

The following command selects only the first six columns of the data frame `df`, creating a new data frame also named `df`.

```
> df <- df[,c(1:6)]
```

2. Selecting based on specific conditions

The following command creates a new data frame called `survived` containing only the rows where the value in the `Survived` column is equal to 1

```
> survived <- subset(df, Survived == 1)
```

The following command creates a new data frame called `notsurvived` containing only the rows where the value in the `Survived` column is equal to 0

```
> notsurvived <- subset(df, Survived == 0)
```

Pivoting Data Towards Specific Questions

Question 1: Determine the quantity of passengers classified under each combination of gender and ticket class.

```
> table(df$Pclass, df$Sex)
```

	female	male
1	94	122
2	76	108
3	144	347

Question 2: Quantify the amount of passengers separated according to their ticket class and survival outcome.

```
> table(df$Pclass, df$Survived)
```

	0	1
1	80	136
2	97	87
3	372	119

Question 3: Estimate the number of passengers divided based on their gender and survival outcome.

```
> table(df$Sex,df$Survived)
```

	0	1
female	81	233
male	468	109

Question 4: Evaluate the number of passengers allocated according to their gender, survival outcome, and ticket class in the Titanic dataset.

```
> table(df$Sex,df$Survived,df$Pclass)
```

```
, , = 1
```

	0	1
female	3	91
male	77	45

```
, , = 2
```


	0	1
female	6	70
male	91	17

```
, , = 3
```

	0	1
female	72	72
male	300	47


Data Visualization

```
> hist(df$Age, main="Distribution of Passengers' Ages", xlab="Age", col="red")
```




This histogram shows us the distribution of ages among the passengers on board the Titanic. We can see that the majority of passengers were between the ages of 20 and 40.

```
> par(mfrow=c(1,2))
> hist(survived$Age, main="Survivors' Age Distribution", xlab="Age",
col="gray")
> hist(notsurvived$Age, main="Non-Survivors' Age Distribution", xlab="Age",
col="red")
```

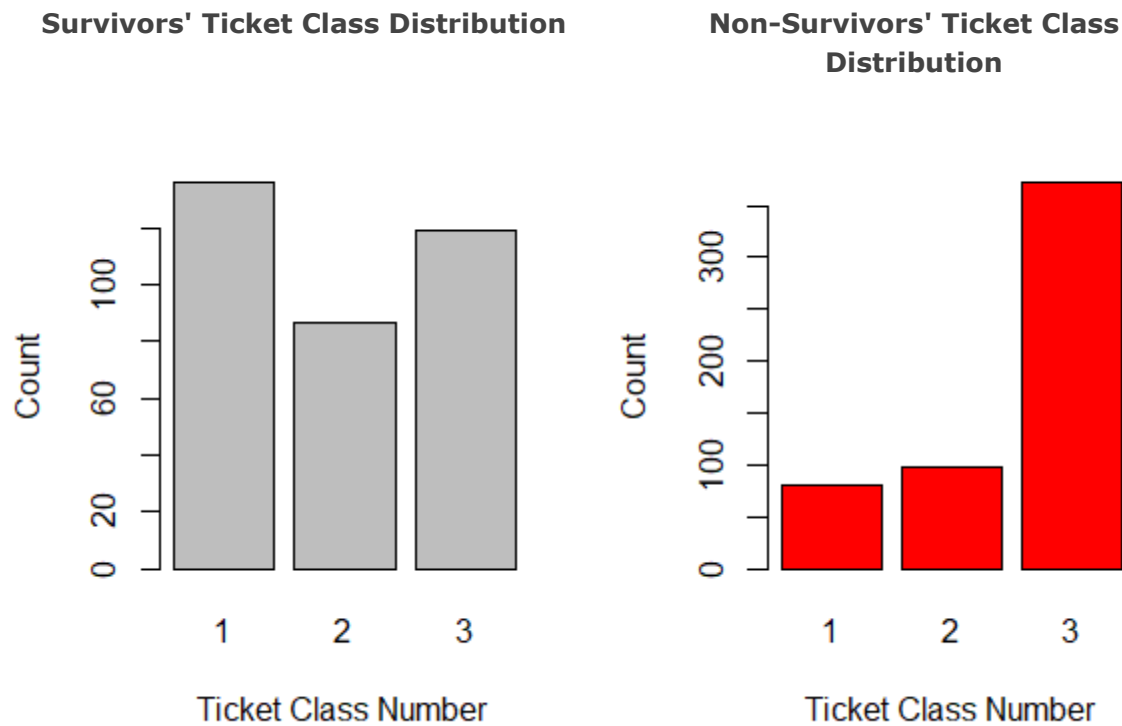


Younger Population Predominant: Majority of the passengers fall under the age range of approximately 20-40 years old, Clear Demarcation Between Survivors and Non-survivors: Survivors generally tend to be younger compared to non-survivors, revealed by the peak in the survivor age distribution occurring notably earlier than the trough in the non-survivor distribution.

```
plot(df$Pclass, main="Distribution of Ticket Classes", xlab="Ticket Class
Number", ylab="Count", col="darkgray")
```



```
> par(mfrow=c(1,2))
> plot(survived$Pclass, main="Survivors' Ticket Class Distribution",
xlab="Ticket Class Number", ylab="Count", col="gray")
> plot(notsurvived$Pclass, main="Non-Survivors' Ticket Class Distribution",
xlab="Ticket Class Number", ylab="Count", col="red")
```



There exists a visible gap between the ticket class distribution of survivors and non-survivors. More survivors belong to 1st and 2nd class compared to non-survivors, implying wealthier passengers had relatively better odds of surviving.